Leveraging AI for Faster Storage Access: a Graph-Neural-Network-Based Prefetcher

Faradawn Yang

Advisor: Haryadi S. Gunawi

Backend storage is slow

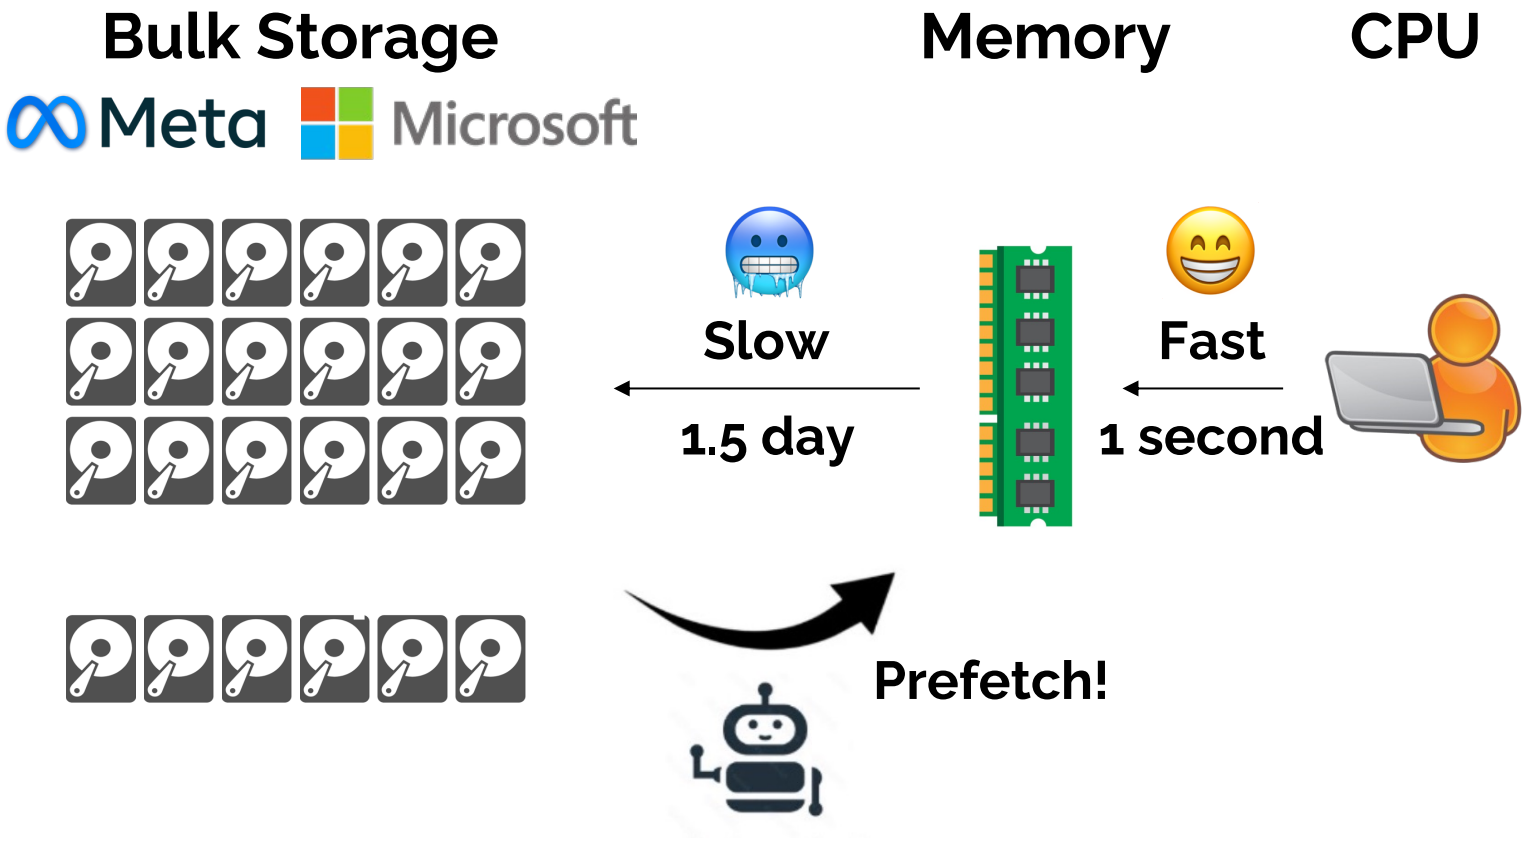

Prefetching improves access speed

No prefetch

With prefetch

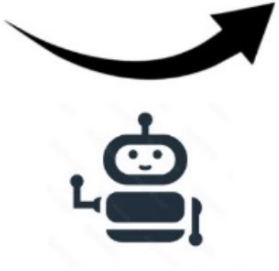

Previous work

['92] Stride

['20] Leap

['21] LSTM

['24] Baleen

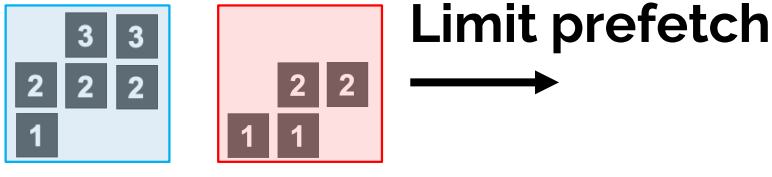

Current problem

Modern workload

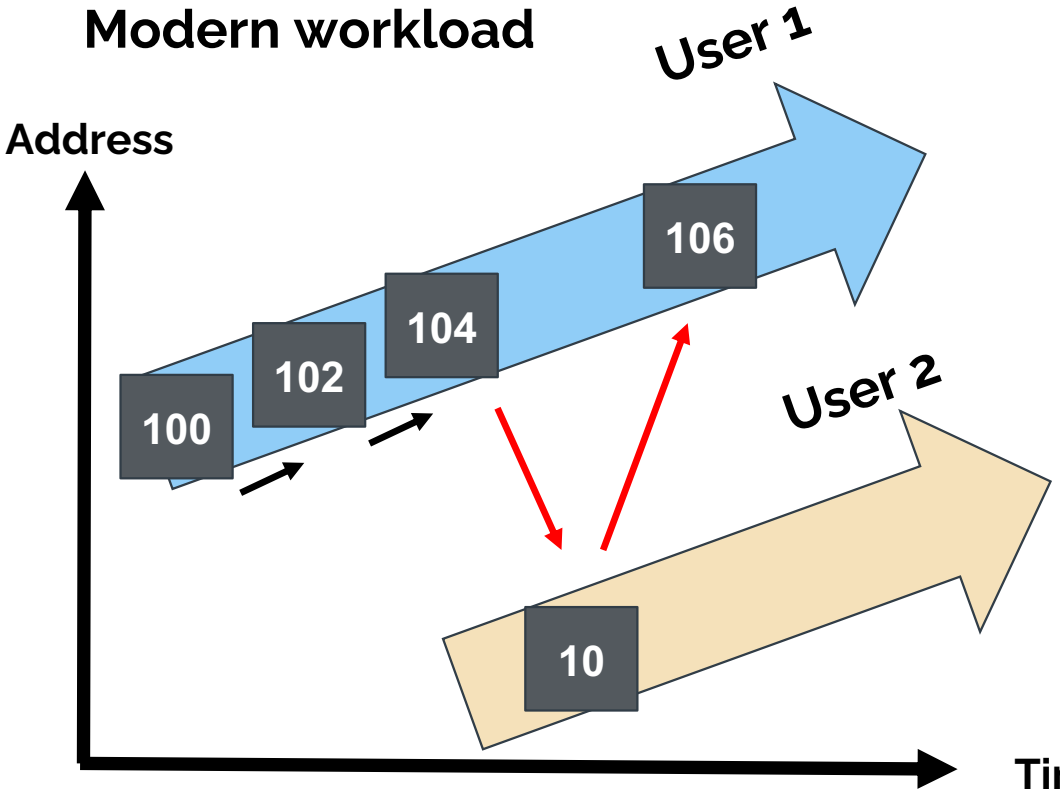

Average hit rate on Alibaba, Tencent, and Microsoft traces

Previous method

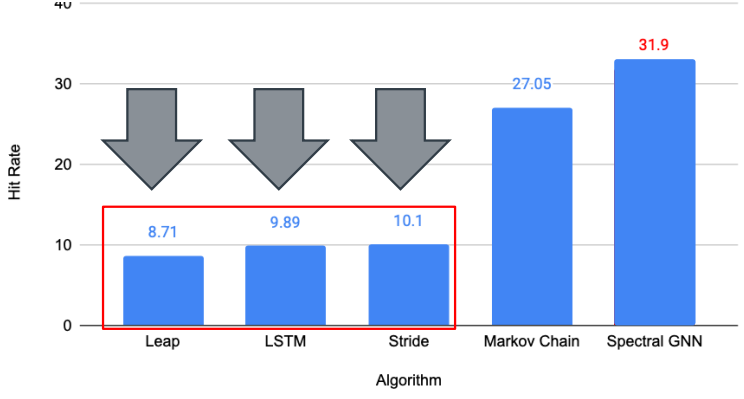

1) Not prone against short-term irregularities.
2) Can't see a global picture.

Methodology

Step 1: Convert to deltas

Address

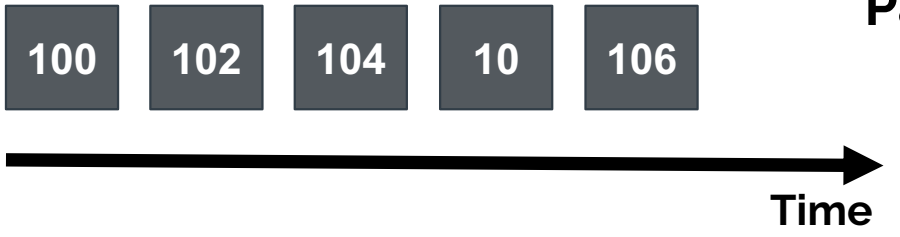

Past 32 accesses.

Delta

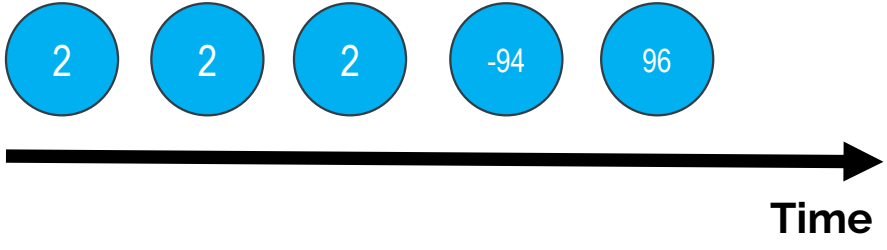

Step 2: Build a graph

Delta

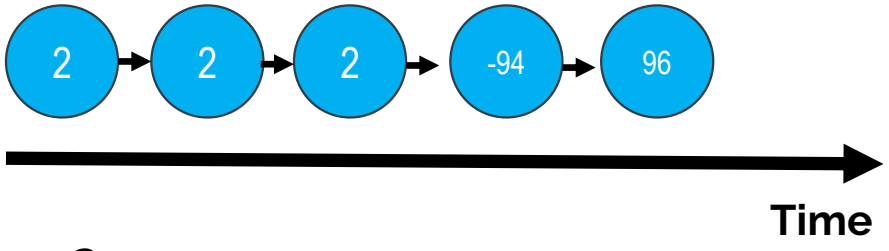

Graph

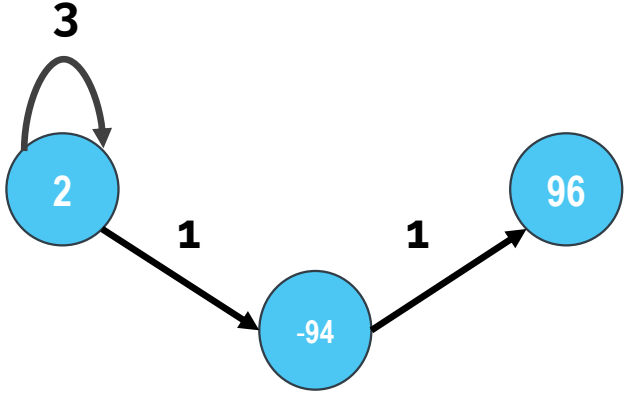

Step 3: Add full connections

Graph

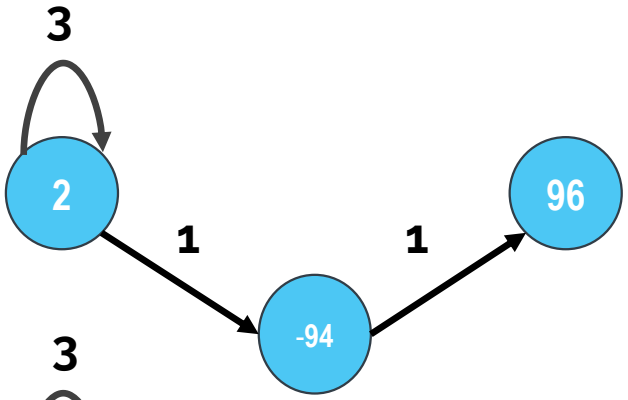

Directed edges capture temporal information.

Fully-connected Graph

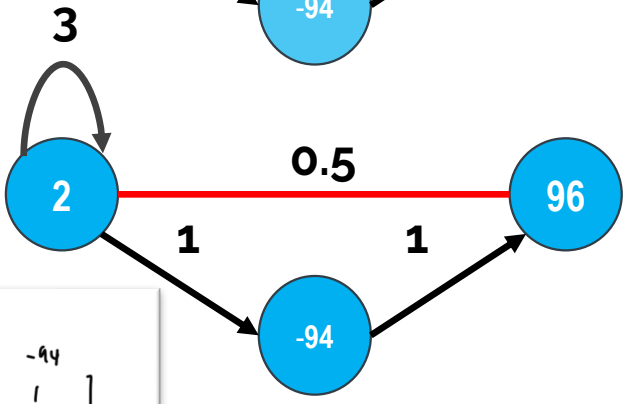

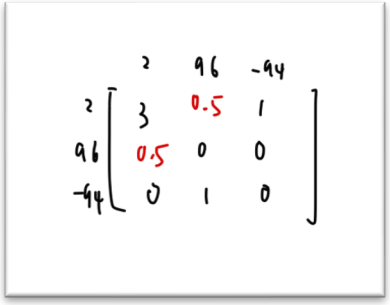

Fully-connected edges reveal spatial pattern.

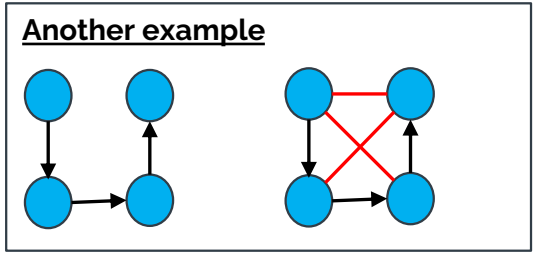

Step 4: Feed into ML

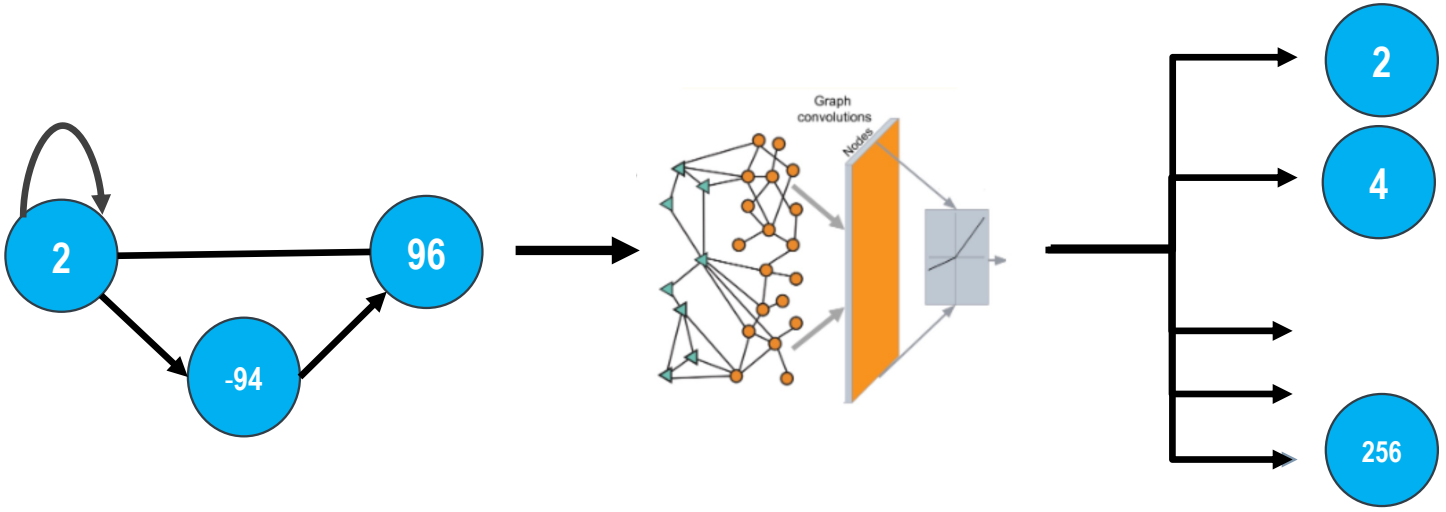

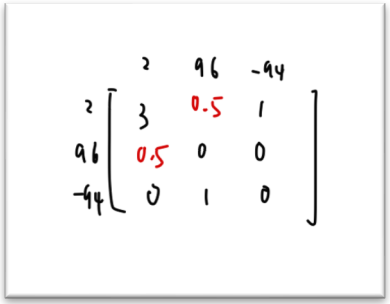

1000 delta addresses.

What ML model

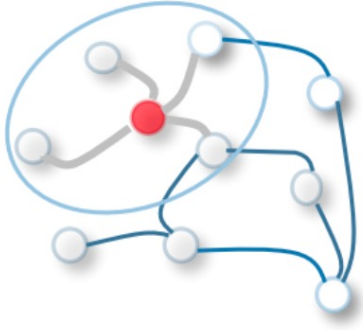

[SGDP '23]

Message passing

Local patterns.
Many iterations.
😞

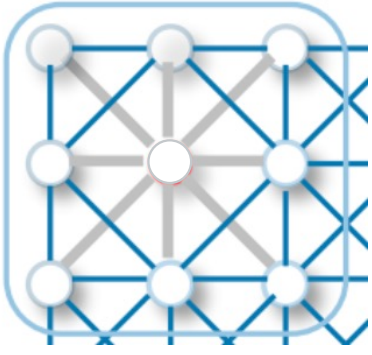

[Ours]

Spectral network

Global structure.
One pass.
😁

Evaluations

Faster access speed

- Setup: python simulator.
- Respected arrival time of requests

Saves 33% access time!

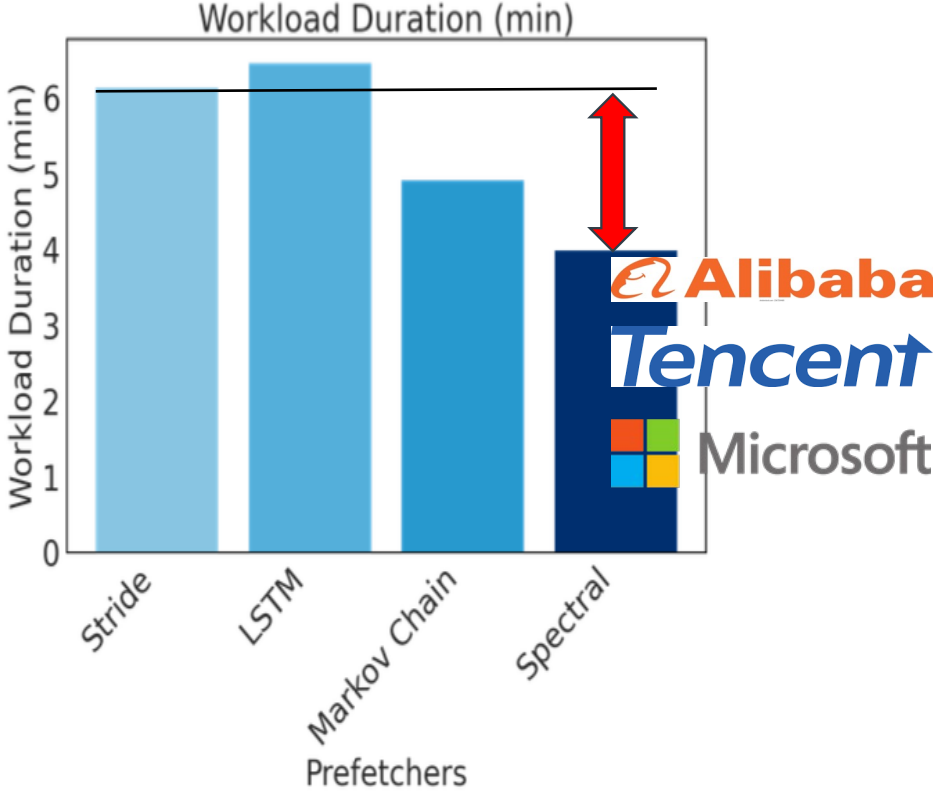

Saves memory and time

- Similar hit rate, but much smaller model and faster training and inference than state-of-the-art, [SGDP '23].

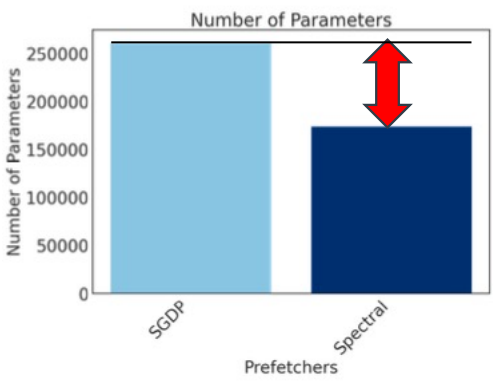

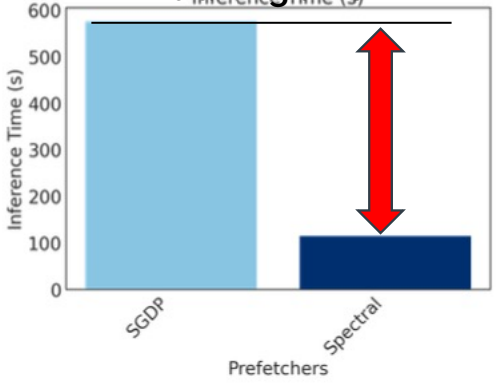

Saves 33.4% model size.

Saves 79% training time.

Conclusion

Entire pipeline

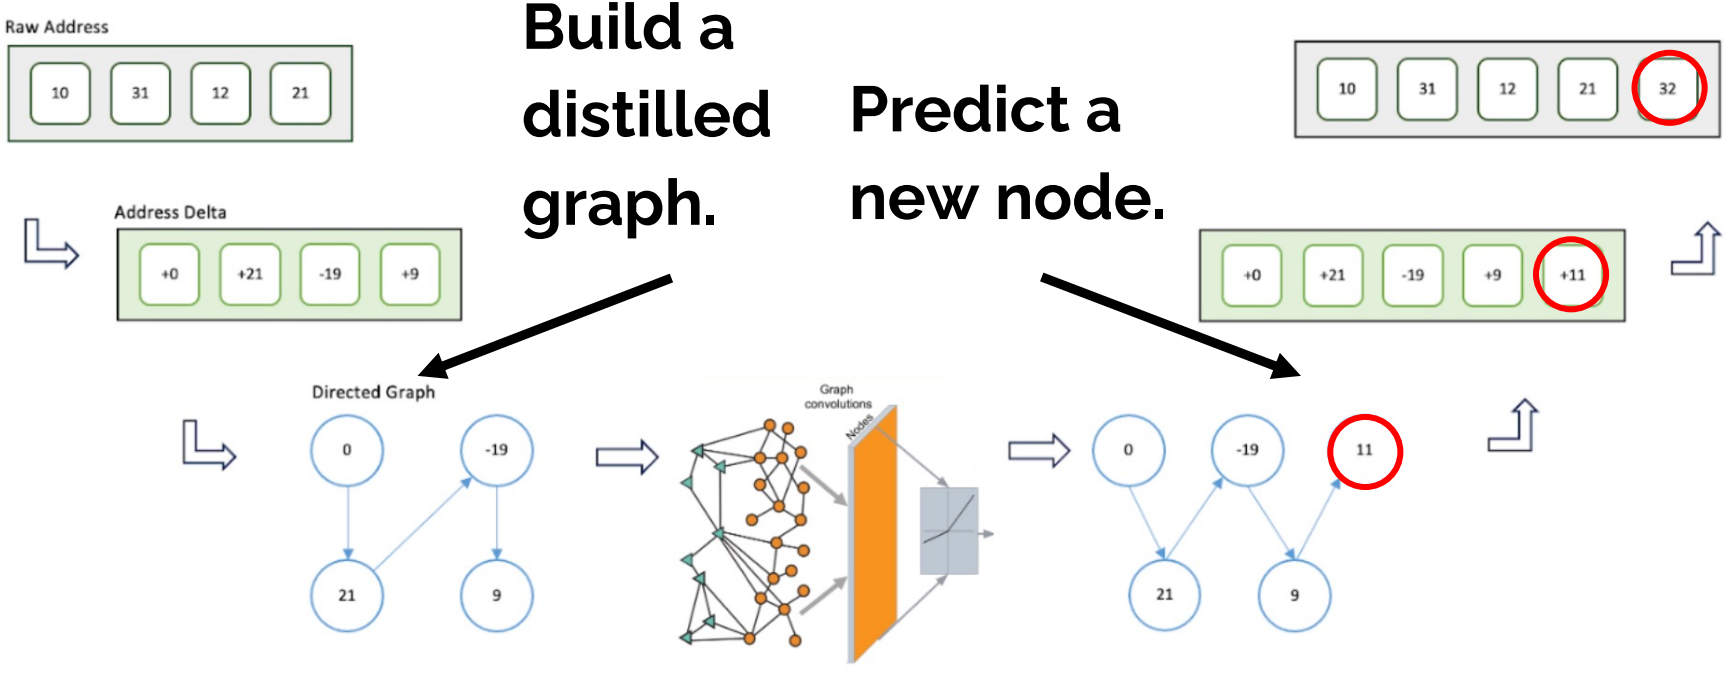

Key takeaway

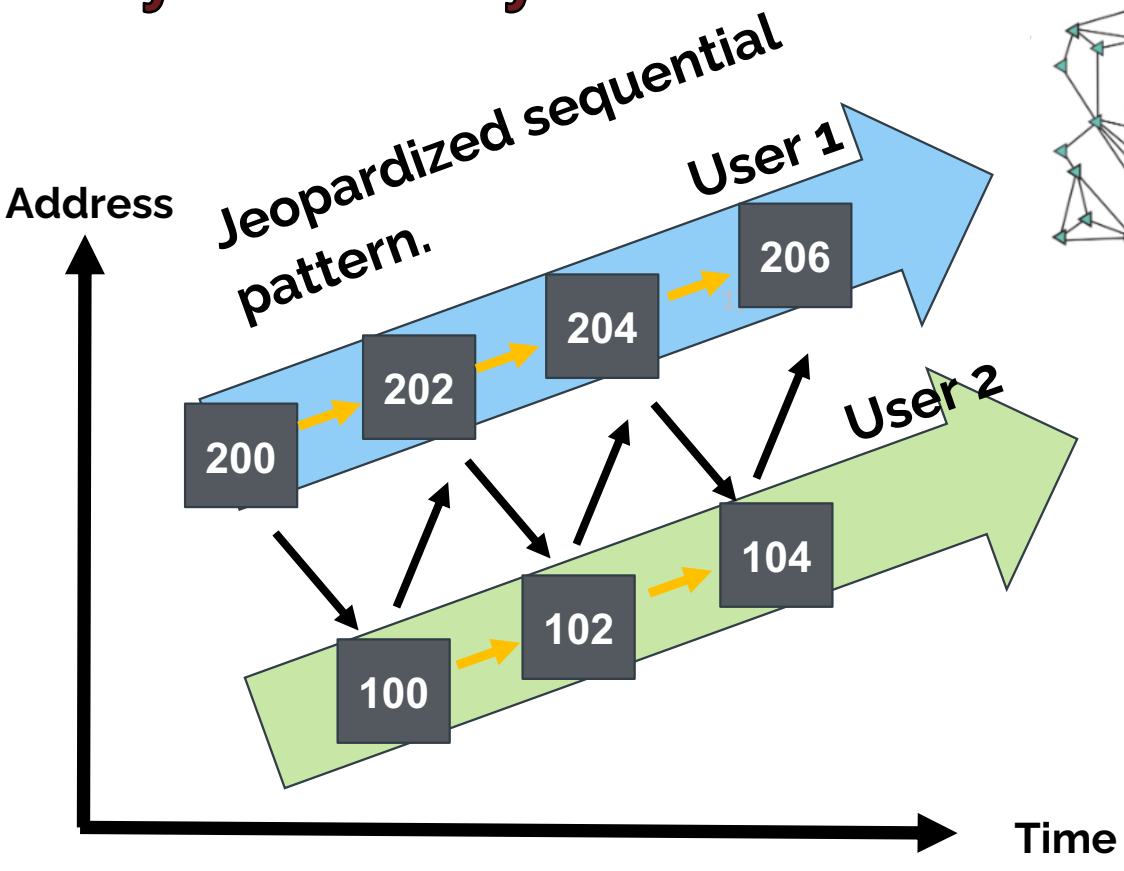

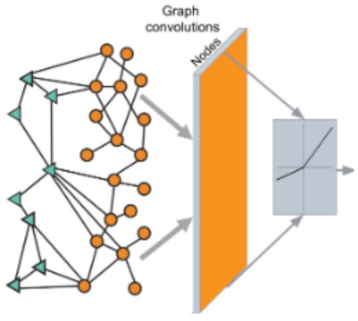

Use graph to detect a global pattern.

Thank you!

Leveraging AI for Faster Storage Access: a Graph-Neural-Network-Based Prefetcher

Backup Slides

Is graph building too costly

- Only putting 32 integers into a matrix.
- Main Bottleneck is disk access (20ms), then inference (2ms), then graph building (1ms).
- Spectral < SGDP < LSTM < transformer.

Higher hit rate

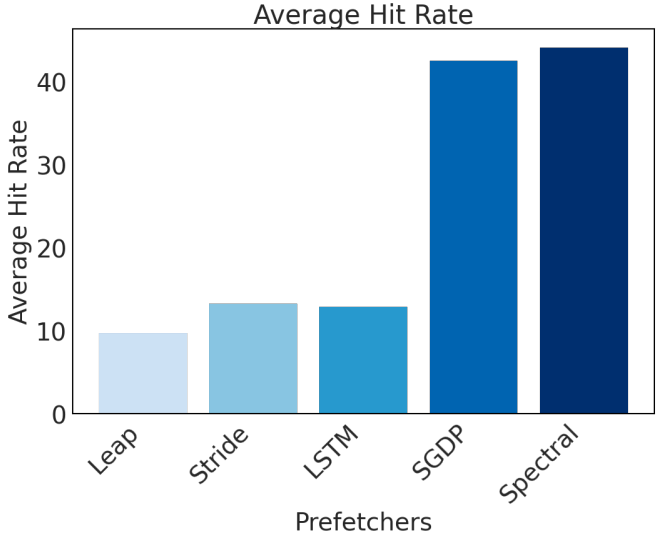

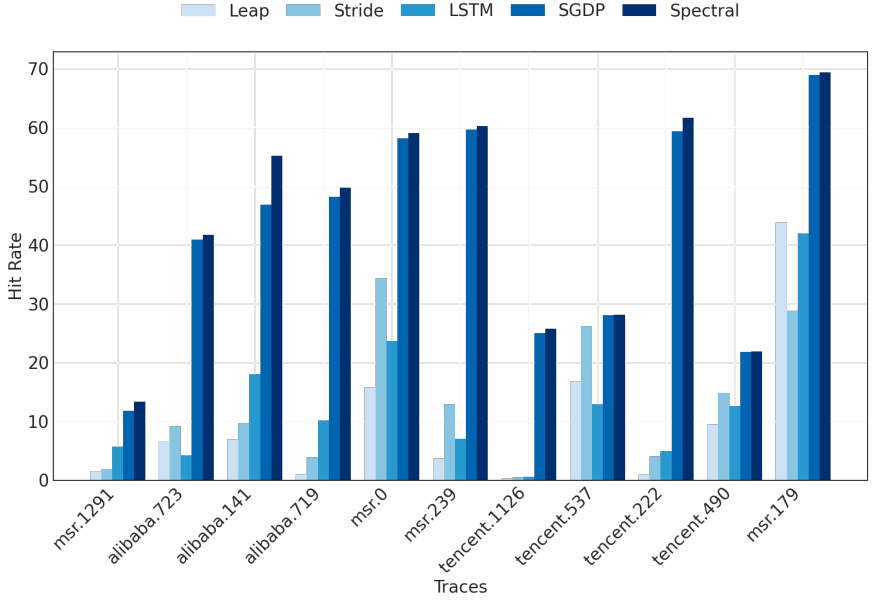

Why spectral method

Message passing

$$m_v^{(t+1)} = \sum_{u \in \mathcal{N}(v)} M^{(t)} \left(h_v^{(t)}, h_u^{(t)}, e_{uv} \right)$$

$$h_v^{(t+1)} = U^{(t)} \left(h_v^{(t)}, m_v^{(t+1)} \right)$$

Many iterations

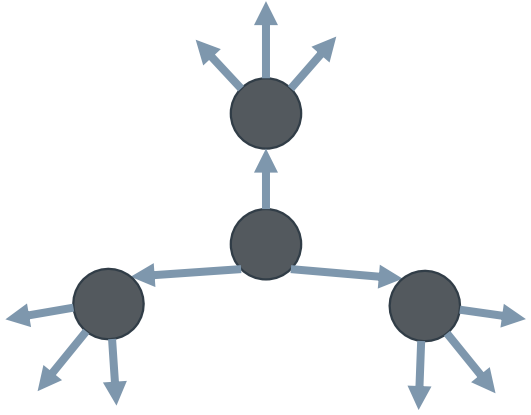

Spectral method

$$H^{(2)} = \sigma_2 \left(\hat{A} \sigma_1 \left(\hat{A} X W^{(0)} \right) W^{(1)} \right)$$

One go